

Hotsuyuki Kawanishi

hotsuyuki.kawanishi@gmail.com | +1 (786) 674-6922 | [linkedin.com/in/hotsuyuki](https://www.linkedin.com/in/hotsuyuki) | github.com/hotsuyuki | hotsuyuki.github.io

SUMMARY

Software Engineer with 4+ years of professional experience in Machine Learning (Large Language Model, Computer Vision) and Autonomous Robot (Localization & Mapping, Perception, Motion Planning). Eligible for US internship under CPT in Summer 2025.

EDUCATION

University of Miami, Coral Gables, Florida | M.S. in **Software Engineering** (GPA: 4.00) **Dec 2025**

Relevant Courses: Machine Learning, Neural Networks, Algorithm Design and Analysis, Software Architecture, Software DevOps

Kanazawa University, Kanazawa, Japan | M.Eng. in **Robotics** (GPA: 3.93) **Mar 2020**

Relevant Courses: Intelligent Robotics, Dynamics and Control, Master's Thesis (on Localization & Mapping for Self-Driving Car)

Chiba University, Chiba, Japan | B.Eng. in **Mechanical Engineering** (GPA: 3.87) **Mar 2017**

Relevant Courses: Linear Algebra, Calculus, Numerical Computation, Analytical Mechanics, Control Theory, Mechatronics

National Institute of Technology Numazu College, Numazu, Japan | A.Eng. in **Computer Science** (GPA: 3.81) **Mar 2015**

Relevant Courses: Data Structures and Algorithms, Programming, Operating Systems, Computer Graphics, Computer Simulation

EXPERIENCE

University of Tokyo AI Lab (Matsuo Lab), Tokyo, Japan **Jan 2024 – Aug 2024**

Machine Learning Engineer, LLM (part-time) | Python, Shell Script, Pre-training, Distributed Systems

- Built an LLM distributed training [pipeline](#) using Microsoft Megatron-DeepSpeed and Hugging Face Transformers, achieving successful training of an OpenAI GPT-based 10B generative model on 160 Nvidia H100 GPUs via Google Cloud GPU cluster.
- Led a lecture on pre-training in an online LLM [course](#) with 4,000+ participants, accomplishing the highest Net Promoter Score.

Woven by Toyota, Tokyo, Japan **May 2021 – Aug 2024**

Software Engineer, Self-Driving Car (full-time) | C++, Python, Prediction, Motion Planning

- Enhanced safety validation for Level 4 driverless autonomous vehicles by designing and implementing a prediction evaluator with Agile methods and software development life cycle best practices, e.g. object-oriented design, code reviews, and testing.
- Decreased driver interventions in JP by 30% by fixing bugs in decision making and trajectory planning with teams in the US, e.g. triaging failures, identifying root causes on a large codebase, articulating solutions, and coordinating vehicle experiments.
- Streamlined and accelerated simulation scenario generation process for the motion planning team by 3x efficiency by creating automation scripts leveraging Applied Intuition API and collaborating with multiple cross-functional teams in the US.

Kanazawa University Advanced Mobility Research Institute, Kanazawa, Japan **Apr 2020 – Mar 2021**

Research Scientist, Self-Driving Car (full-time) | C++, Python, Localization & Mapping, Perception

- Improved map data collection process for autonomous vehicles by 2x efficiency by creating a real-time map generation tool.
- Developed perception features by applying research theories to practical applications in a fast-paced environment, including a speed bump detector using LiDAR point cloud and probabilistic methods (successfully delivered to a customer company) and a lane line visibility checker using LiDAR-camera sensor fusion and linear algebra (published in an international journal [paper](#)).
- Successfully mentored 4 lab students, resulting in 1 [patent](#) filing and 1 [paper](#) publication, by providing hands-on guidance.

Cultural Vistas, New York City, New York **Jul 2017 – Dec 2017**

Software Engineer, Web (internship) | JavaScript, PHP, SQL, Relational Database, Full-stack

- Developed and maintained the web front-end and back-end in a diverse, multicultural environment of 5 software engineers.

SKILLS

Programming Languages: C++, Python, C, C#, JavaScript, TypeScript, PHP, SQL, HTML, CSS, Shell Script

Frameworks / Technologies: PyTorch, TensorFlow, Keras, DeepSpeed, Hugging Face Hub, Transformers, W&B, CUDA, ROS, Jupyter Notebook, NumPy, Matplotlib, Unity, Niantic ARDK, React, FastAPI, AWS, Azure, GCP, Docker, Jenkins, Bazel, Git, Linux

PROJECTS

[LLM Distributed Training Pipeline](#) | Python, Shell Script, PyTorch, DeepSpeed, Transformers **Jan 2024 – Apr 2024**

- Scalable distributed training pipeline for LLMs on GPU clusters using DeepSpeed ZeRO, 3D Parallelism, and Transformers.

[HuggingGPT by GPT Function Calling](#) | TypeScript, React, Python, FastAPI, Azure, OpenAI API **Jul 2023 – Jul 2023**

- Agentic AI web app to interact with various Hugging Face APIs through natural language by utilizing GPT Function Calling.

[C++ Deep Learning Framework from Scratch](#) | C++, Object-Oriented Programming, Unit Testing **Jun 2022 – Aug 2022**

- PyTorch-like deep learning framework architected and implemented from scratch in C++, e.g. layers, optimizers, and autograd.

[YOLOv5 Object Detection Model in PyTorch C++](#) | C++, PyTorch, Convolutional Neural Network **Dec 2020 – Dec 2020**

- YOLOv5 implementation using PyTorch C++ API (LibTorch) instead of Python for efficient integration in C++ applications.

[Robotics Algorithms in Jupyter Notebook](#) | Python, Robotics, SLAM, Localization, Path Planning **Aug 2020 – Sep 2020**

- Python codes for robotics algorithms, e.g. Simultaneous Localization & Mapping (SLAM), localization, and path planning.